

Assignment 3

All questions are weighted the same in this assignment. This assignment requires more individual learning than the last one did - you are encouraged to check out the [pandas documentation](http://pandas.pydata.org/pandas-docs/stable/) (<http://pandas.pydata.org/pandas-docs/stable/>) to find functions or methods you might not have used yet, or ask questions on [Stack Overflow](http://stackoverflow.com/) (<http://stackoverflow.com/>) and tag them as pandas and python related. All questions are worth the same number of points except question 1 which is worth 17% of the assignment grade.

Note: Questions 2-13 rely on your question 1 answer.

```
In [1]: # Import Libraries
import sys
import pandas as pd
import numpy as np

%matplotlib inline
import matplotlib.pyplot as plt
plt.style.use('seaborn-whitegrid')

import warnings
warnings.filterwarnings('ignore')

print('You\'re running python %s' % sys.version.split(' ')[0])
```

You're running python 3.7.3

Question 1

Load the energy data from the file `assets/Energy Indicators.xls`, which is a list of indicators of [energy supply and renewable electricity production](http://unstats.un.org/unsd/environment/excel_file_tables/2013/Energy%20Indicators.xls) ([assets/Energy%20Indicators.xls](http://unstats.un.org/unsd/environment/excel_file_tables/2013/Energy%20Indicators.xls)) from the [United Nations](http://unstats.un.org/unsd/environment/excel_file_tables/2013/Energy%20Indicators.xls) (http://unstats.un.org/unsd/environment/excel_file_tables/2013/Energy%20Indicators.xls) for the year 2013, and should be put into a DataFrame with the variable name of **Energy**.

Keep in mind that this is an Excel file, and not a comma separated values file. Also, make sure to exclude the footer and header information from the datafile. The first two columns are unnecessary, so you should get rid of them, and you should change the column labels so that the columns are:

```
['Country', 'Energy Supply', 'Energy Supply per Capita', '% Renewable']
```

Convert Energy Supply to gigajoules (**Note: there are 1,000,000 gigajoules in a petajoule**). For all countries which have missing data (e.g. data with "...") make sure this is reflected as `np.NaN` values.

Rename the following list of countries (for use in later questions):

```
"Republic of Korea": "South Korea",
"United States of America": "United States",
"United Kingdom of Great Britain and Northern Ireland": "United Kingdom",
"China, Hong Kong Special Administrative Region": "Hong Kong"
```

There are also several countries with numbers and/or parenthesis in their name. Be sure to remove these, e.g. `'Bolivia (Plurinational State of)'` should be `'Bolivia'`. `'Switzerland17'` should be `'Switzerland'`.

Next, load the GDP data from the file `assets/world_bank.csv`, which is a csv containing countries' GDP from 1960 to 2015 from [World Bank](http://data.worldbank.org/indicator/NY.GDP.MKTP.CD) (<http://data.worldbank.org/indicator/NY.GDP.MKTP.CD>). Call this DataFrame **GDP**.

Make sure to skip the header, and rename the following list of countries:

```
"Korea, Rep.": "South Korea",
"Iran, Islamic Rep.": "Iran",
"Hong Kong SAR, China": "Hong Kong"
```

Finally, load the [Scimago Journal and Country Rank data for Energy Engineering and Power Technology](http://www.scimagojr.com/countryrank.php?category=2102) (<http://www.scimagojr.com/countryrank.php?category=2102>) from the file `assets/scimagojr-3.xlsx`, which ranks countries based on their journal contributions in the aforementioned area. Call this DataFrame **ScimEn**.

Join the three datasets: GDP, Energy, and ScimEn into a new dataset (using the intersection of country names). Use only the last 10 years (2006-2015) of GDP data and only the top 15 countries by Scimagojr 'Rank' (Rank 1 through 15).

The index of this DataFrame should be the name of the country, and the columns should be ['Rank', 'Documents', 'Citable documents', 'Citations', 'Self-citations', 'Citations per document', 'H index', 'Energy Supply', 'Energy Supply per Capita', '% Renewable', '2006', '2007', '2008', '2009', '2010', '2011', '2012', '2013', '2014', '2015'].

This function should return a DataFrame with 20 columns and 15 entries, and the rows of the DataFrame should be sorted by "Rank".

```

In [2]: '''
@author: Steven Ponce
Date: 22 April 2021
'''

def answer_one():
    # YOUR CODE HERE
    #raise NotImplementedError()

    # Loading the dataset and save DataFrame with the variable name of Energy
    # Dropping the first two columns and renaming the remaining columns
    Energy = pd.read_excel('assets/Energy Indicators.xls',
                          na_values=["..."],
                          header = None, skiprows=18, skipfooter= 38,
                          usecols=[2,3,4,5],
                          names=['Country', 'Energy Supply', 'Energy Supply per Capita', '% Renewable'])

    # Convert Energy Supply to gigajoules (Note: there are 1,000,000 gigajoules in a petajoule).
    Energy['Energy Supply'] = Energy['Energy Supply'].apply(lambda x: x * 1000000)

    # There are also several countries with numbers and/or parenthesis in their name. Be sure to remove them.
    Energy['Country'] = Energy['Country'].str.replace(r" \(.*\)", "")
    Energy['Country'] = Energy['Country'].str.replace(r"\d*", "")

    # Rename the following list of countries:
    Energy['Country'] = Energy['Country'].replace({'Republic of Korea' : 'South Korea',
                                                  'United States of America' : 'United States',
                                                  'United Kingdom of Great Britain and Northern Ireland': 'United Kingdom',
                                                  'China, Hong Kong Special Administrative Region': 'Hong Kong'})

    # Next, load the GDP data from the file assets/world_bank.csv. Call this DataFrame GDP.
    GDP = pd.read_csv('assets/world_bank.csv', skiprows=4)

    # Rename the following list of countries:
    GDP['Country Name'] = GDP['Country Name'].replace({'Korea, Rep.': 'South Korea',
                                                       'Iran, Islamic Rep.': 'Iran',
                                                       'Hong Kong SAR, China' : 'Hong Kong'})

    # Finally, load the file assets/scimagojr-3.xlsx, which ranks countries based on their journal contributions in
    # the aforementioned area. Call this DataFrame ScimEn.
    ScimEn = pd.read_excel('assets/scimagojr-3.xlsx')

    """
    Join the three datasets: GDP, Energy, and ScimEn into a new dataset (using the intersection of country names).

    Use only the last 10 years (2006-2015) of GDP data and only the top 15 countries by Scimagojr 'Rank' (Rank 1
    through 15).

    The index of this DataFrame should be the name of the country, and the columns should be ['Rank', 'Documents',
    'Citable documents', 'Citations', 'Self-citations', 'Citations per document', 'H index', 'Energy Supply',
    'Energy Supply per Capita', '% Renewable', '2006', '2007', '2008', '2009', '2010', '2011', '2012', '2013',
    '2014', '2015'].

    This function should return a DataFrame with 20 columns and 15 entries, and the rows of the DataFrame
    should be sorted by "Rank".

    OUTPUT: This function should return a DataFrame with 20 columns and 15 entries, and the rows of the
    DataFrame should be sorted by "Rank".
    """

    # Join the three datasets: GDP, Energy, and ScimEn into a new dataset
    Energy_ScimEn = pd.merge(ScimEn, Energy, how='inner', left_on='Country', right_on='Country')
    Energy_ScimEn = Energy_ScimEn[Energy_ScimEn['Rank'] <= 15]

    GDP.rename(columns = {'Country Name': 'Country'}, inplace=True)
    GDP = GDP.loc[:, ['2006', '2007', '2008', '2009', '2010', '2011', '2012', '2013', '2014', '2015', 'Country']]

    Energy_ScimEn_GDP = pd.merge(Energy_ScimEn, GDP, how="inner", left_on="Country", right_on="Country").set_index('Country')

    return Energy_ScimEn_GDP

```

In [3]: answer_one().shape

Out[3]: (15, 20)

```
In [4]: answer_one()
```

Out[4]:

	Rank	Documents	Citable documents	Citations	Self-citations	Citations per document	H index	Energy Supply	Energy Supply per Capita	% Renewable	2006	2007	
Country													
China	1	127050	126767	597237	411683	4.70	138	1.271910e+11	93.0	19.754910	3.992331e+12	4.559041e+12	4.9977
United States	2	96661	94747	792274	265436	8.20	230	9.083800e+10	286.0	11.570980	1.479230e+13	1.505540e+13	1.5011
Japan	3	30504	30287	223024	61554	7.31	134	1.898400e+10	149.0	10.232820	5.496542e+12	5.617036e+12	5.5585
United Kingdom	4	20944	20357	206091	37874	9.84	139	7.920000e+09	124.0	10.600470	2.419631e+12	2.482203e+12	2.4706
Russian Federation	5	18534	18301	34266	12422	1.85	57	3.070900e+10	214.0	17.288680	1.385793e+12	1.504071e+12	1.5830
Canada	6	17899	17620	215003	40930	12.01	149	1.043100e+10	296.0	61.945430	1.564469e+12	1.596740e+12	1.6127
Germany	7	17027	16831	140566	27426	8.26	126	1.326100e+10	165.0	17.901530	3.332891e+12	3.441561e+12	3.4788
India	8	15005	14841	128763	37209	8.58	115	3.319500e+10	26.0	14.969080	1.265894e+12	1.374865e+12	1.4283
France	9	13153	12973	130632	28601	9.93	114	1.059700e+10	166.0	17.020280	2.607840e+12	2.669424e+12	2.6746
South Korea	10	11983	11923	114675	22595	9.57	104	1.100700e+10	221.0	2.279353	9.410199e+11	9.924316e+11	1.0205
Italy	11	10964	10794	111850	26661	10.20	106	6.530000e+09	109.0	33.667230	2.202170e+12	2.234627e+12	2.2111
Spain	12	9428	9330	123336	23964	13.08	115	4.923000e+09	106.0	37.968590	1.414823e+12	1.468146e+12	1.4845
Iran	13	8896	8819	57470	19125	6.46	72	9.172000e+09	119.0	5.707721	3.895523e+11	4.250646e+11	4.2899
Australia	14	8831	8725	90765	15606	10.28	107	5.386000e+09	231.0	11.810810	1.021939e+12	1.060340e+12	1.0996
Brazil	15	8668	8596	60702	14396	7.00	86	1.214900e+10	59.0	69.648030	1.845080e+12	1.957118e+12	2.0568

```
In [5]: assert type(answer_one()) == pd.DataFrame, "Q1: You should return a DataFrame!"

assert answer_one().shape == (15,20), "Q1: Your DataFrame should have 20 columns and 15 entries!"
```


```
In [6]: # Cell for autograder.
```

Question 2

The previous question joined three datasets then reduced this to just the top 15 entries. When you joined the datasets, but before you reduced this to the top 15 items, how many entries did you lose?

This function should return a single number.

```
In [7]: %HTML
<svg width="800" height="300">
  <circle cx="150" cy="180" r="80" fill-opacity="0.2" stroke="black" stroke-width="2" fill="blue" />
  <circle cx="200" cy="100" r="80" fill-opacity="0.2" stroke="black" stroke-width="2" fill="red" />
  <circle cx="100" cy="100" r="80" fill-opacity="0.2" stroke="black" stroke-width="2" fill="green" />
  <line x1="150" y1="125" x2="300" y2="150" stroke="black" stroke-width="2" fill="black" stroke-dasharray="5,3"/>
  <text x="300" y="165" font-family="Verdana" font-size="35">Everything but this!</text>
</svg>
```



Everything but this!

```

In [8]: '''
@author: Steven Ponce
Date: 22 April 2021
'''

def answer_two():
    # YOUR CODE HERE
    #raise NotImplementedError()

    # Loading the dataset and save DataFrame with the variable name of Energy
    # Dropping the first two columns and renaming the remaining columns
    Energy = pd.read_excel('assets/Energy Indicators.xls',
                           na_values=["..."], header = None, skiprows=18, skipfooter= 38,
                           usecols=[2,3,4,5],
                           names=['Country', 'Energy Supply', 'Energy Supply per Capita', '% Renewable'])

    # Convert Energy Supply to gigajoules (Note: there are 1,000,000 gigajoules in a petajoule).
    Energy['Energy Supply'] = Energy['Energy Supply'].apply(lambda x: x * 1000000)

    # There are also several countries with numbers and/or parenthesis in their name. Be sure to remove them.
    Energy['Country'] = Energy['Country'].str.replace(r" \(.*\)", "")
    Energy['Country'] = Energy['Country'].str.replace(r"\d*", "")

    # Rename the following list of countries:
    Energy['Country'] = Energy['Country'].replace({'Republic of Korea' : 'South Korea',
                                                  'United States of America' : 'United States',
                                                  'United Kingdom of Great Britain and Northern Ireland': 'United Kingdom',
                                                  'China, Hong Kong Special Administrative Region': 'Hong Kong'})

    # Next, Load the GDP data from the file assets/world_bank.csv. Call this DataFrame GDP.
    GDP = pd.read_csv('assets/world_bank.csv', skiprows=4)

    # Rename the following list of countries:
    GDP['Country Name'] = GDP['Country Name'].replace({'Korea, Rep.': 'South Korea',
                                                       'Iran, Islamic Rep.': 'Iran',
                                                       'Hong Kong SAR, China' : 'Hong Kong'})

    # Finally, Load the file assets/scimagojr-3.xlsx, which ranks countries based on their journal contributions in
    # the aforementioned area. Call this DataFrame ScimEn.
    ScimEn = pd.read_excel('assets/scimagojr-3.xlsx')

    # inner1 = Energy_ScimEn_1
    # inner2 = Energy_ScimEn_GDP_1
    # outer1 = Energy_ScimEn_2
    # outer2 = Energy_ScimEn_GDP_2

    # inner join
    Energy_ScimEn_1 = pd.merge(ScimEn, Energy, how='inner', left_on='Country', right_on='Country')

    GDP.rename(columns = {'Country Name': 'Country'}, inplace=True)
    GDP = GDP.loc[:, ['2006', '2007', '2008', '2009', '2010', '2011', '2012', '2013', '2014', '2015', 'Country']]

    Energy_ScimEn_GDP_1 = pd.merge(Energy_ScimEn_1, GDP, how='inner', left_on='Country',
                                   right_on='Country').set_index('Country')

    # outer join
    Energy_ScimEn_2 = pd.merge(ScimEn, Energy, how='outer', left_on='Country', right_on='Country')
    Energy_ScimEn_GDP_2 = pd.merge(Energy_ScimEn_2, GDP, how='outer', left_on='Country',
                                   right_on='Country').set_index('Country')

    # outer minus join
    return len(Energy_ScimEn_GDP_2) - len(Energy_ScimEn_GDP_1);

```

In [9]: answer_two()

Out[9]: 156

In [10]: assert type(answer_two()) == int, "Q2: You should return an int number!"

Question 3

What are the top 15 countries for average GDP over the last 10 years?

This function should return a Series named `avgGDP` with 15 countries and their average GDP sorted in descending order.

```
In [11]: '''
    @author: Steven Ponce
    Date: 22 April 2021
    '''

    def answer_three():
        # YOUR CODE HERE
        #raise NotImplementedError()

        df = answer_one()
        avgGDP = df[['2006', '2007', '2008', '2009', '2010', '2011', '2012',
                    '2013', '2014', '2015']].mean(axis=1).rename('aveGDP').sort_values(ascending=False)

        return avgGDP
```

```
In [12]: answer_three().shape
```

Out[12]: (15,)

```
In [13]: answer_three()
```

Out[13]:

Country	
United States	1.536434e+13
China	6.348609e+12
Japan	5.542208e+12
Germany	3.493025e+12
France	2.681725e+12
United Kingdom	2.487907e+12
Brazil	2.189794e+12
Italy	2.120175e+12
India	1.769297e+12
Canada	1.660647e+12
Russian Federation	1.565459e+12
Spain	1.418078e+12
Australia	1.164043e+12
South Korea	1.106715e+12
Iran	4.441558e+11

Name: aveGDP, dtype: float64

```
In [14]: assert type(answer_three()) == pd.Series, "Q3: You should return a Series!"
```

Question 4

By how much had the GDP changed over the 10 year span for the country with the 6th largest average GDP?

This function should return a single number.

```
In [15]: '''
    @author: Steven Ponce
    Date: 22 April 2021
    '''

    def answer_four():
        # YOUR CODE HERE
        #raise NotImplementedError()

        df = answer_one()
        avgGDP = df[['2006', '2007', '2008', '2009', '2010', '2011', '2012',
                    '2013', '2014', '2015']].mean(axis=1).rename('avgGDP').sort_values(ascending=False)

        sixth_largest_avgGDP = df.iloc[3]['2015'] - df.iloc[3]['2006']

        return sixth_largest_avgGDP
```

```
In [16]: answer_four()
```

Out[16]: 246702696075.3999

```
In [17]: # Cell for autograder.
```

Question 5

What is the mean energy supply per capita?

This function should return a single number.

```
In [18]: '''
@author: Steven Ponce
Date: 22 April 2021
'''

def answer_five():
    # YOUR CODE HERE
    #raise NotImplementedError()

    df = answer_one()
    mean_energy_per_capita = df['Energy Supply per Capita'].mean()

    return mean_energy_per_capita
```

```
In [19]: answer_five()
```

```
Out[19]: 157.6
```

```
In [20]: # Cell for autograder.
```

Question 6

What country has the maximum % Renewable and what is the percentage?

This function should return a tuple with the name of the country and the percentage.

```
In [21]: '''
@author: Steven Ponce
Date: 22 April 2021
'''

def answer_six():
    # YOUR CODE HERE
    #raise NotImplementedError()

    df = answer_one()
    country_max_renewable = df['% Renewable'].idxmax(), df['% Renewable'].max()

    return country_max_renewable
```

```
In [22]: answer_six()
```

```
Out[22]: ('Brazil', 69.64803)
```

```
In [23]: assert type(answer_six()) == tuple, "Q6: You should return a tuple!"

assert type(answer_six()[0]) == str, "Q6: The first element in your result should be the name of the country!"
```

Question 7

Create a new column that is the ratio of Self-Citations to Total Citations. What is the maximum value for this new column, and what country has the highest ratio?

This function should return a tuple with the name of the country and the ratio.

```
In [24]: '''
@author: Steven Ponce
Date: 23 April 2021
'''

def answer_seven():
    # YOUR CODE HERE
    #raise NotImplementedError()

    df = answer_one()
    df['Citation Ratio'] = df['Self-citations'] / df['Citations']
    max_citation_ratio = df['Citation Ratio'].idxmax(), df['Citation Ratio'].max()

    return max_citation_ratio
```

```
In [25]: answer_seven()
```

```
Out[25]: ('China', 0.6893126179389422)
```

```
In [26]: assert type(answer_seven()) == tuple, "Q7: You should return a tuple!"

assert type(answer_seven()[0]) == str, "Q7: The first element in your result should be the name of the country!"
```

Question 8

Create a column that estimates the population using Energy Supply and Energy Supply per capita. What is the third most populous country according to this estimate?

This function should return the name of the country

```
In [27]: '''
@author: Steven Ponce
Date: 23 April 2021
'''

def answer_eight():
    # YOUR CODE HERE
    #raise NotImplementedError()

    df = answer_one()
    df['Population Estimate'] = df['Energy Supply'] / df['Energy Supply per Capita']
    third_populous_country = df.sort_values(by='Population Estimate', ascending=False).iloc[2].name

    return third_populous_country
```

```
In [28]: answer_eight()
```

```
Out[28]: 'United States'
```

```
In [29]: assert type(answer_eight()) == str, "Q8: You should return the name of the country!"
```

Question 9

Create a column that estimates the number of citable documents per person. What is the correlation between the number of citable documents per capita and the energy supply per capita? Use the `.corr()` method, (Pearson's correlation).

This function should return a single number.

(Optional: Use the built-in function `plot9()` to visualize the relationship between Energy Supply per Capita vs. Citable docs per Capita)

```
In [30]: '''
@author: Steven Ponce
Date: 23 April 2021
'''

def answer_nine():
    # YOUR CODE HERE
    # raise NotImplementedError()

    df = answer_one()
    df = df.reset_index()
    df['Population Estimate'] = df['Energy Supply'] / df['Energy Supply per Capita']
    df['Citable docs per Capita'] = df['Citable documents'] / df['Population Estimate']

    correlation = df[['Energy Supply per Capita',
                      'Citable docs per Capita']].corr().ix['Energy Supply per Capita', 'Citable docs per Capita']

    return correlation
```

```
In [31]: answer_nine()

Out[31]: 0.7940010435442943
```

```
In [32]: '''
@author: Steven Ponce
Date: 23 April 2021
'''

def plot_nine():

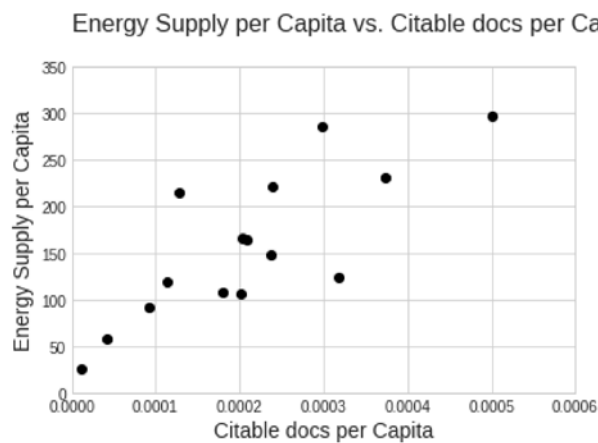
    df = answer_one()
    df['Population Estimate'] = df['Energy Supply'] / df['Energy Supply per Capita']
    df['Citable docs per Capita'] = df['Citable documents'] / df['Population Estimate']

    x_data = df['Citable docs per Capita']
    y_data = df['Energy Supply per Capita']

    plt.scatter(x_data, y_data, marker = 'o', color='black')

    title=('Energy Supply per Capita vs. Citable docs per Capita\n')
    plt.title(title, loc='left', fontsize=16)
    plt.xlabel('Citable docs per Capita', fontsize=14)
    plt.ylabel('Energy Supply per Capita', fontsize=14)
    plt.ylim(0, 350)
    plt.xlim(0.0000, 0.0006)
```

```
In [33]: plot_nine()
```



```
In [34]: assert answer_nine() >= -1. and answer_nine() <= 1., "Q9: A valid correlation should between -1 to 1!"
```

Question 10

Create a new column with a 1 if the country's % Renewable value is at or above the median for all countries in the top 15, and a 0 if the country's % Renewable value is below the median.

This function should return a series named `HighRenew` whose index is the country name sorted in ascending order of rank.


```
In [35]: '''
@author: Steven Ponce
Date: 23 April 2021
'''

def answer_ten():
    # YOUR CODE HERE
    #raise NotImplementedError()

    df = answer_one()
    Median = df['% Renewable'].median()
    HighRenew = df['HighRenew'] = df['% Renewable'].apply(lambda x:0 if x<Median else 1 )

    return HighRenew
```

```
In [36]: answer_ten()
```

```
Out[36]: Country
China          1
United States  0
Japan          0
United Kingdom 0
Russian Federation 1
Canada         1
Germany        1
India          0
France         1
South Korea    0
Italy          1
Spain         1
Iran          0
Australia     0
Brazil        1
Name: % Renewable, dtype: int64
```

```
In [37]: assert type(answer_ten()) == pd.Series, "Q10: You should return a Series!"
```

Question 11

Use the following dictionary to group the Countries by Continent, then create a DataFrame that displays the sample size (the number of countries in each continent bin), and the sum, mean, and std deviation for the estimated population of each country.

```
ContinentDict = {'China':'Asia',
                 'United States':'North America',
                 'Japan':'Asia',
                 'United Kingdom':'Europe',
                 'Russian Federation':'Europe',
                 'Canada':'North America',
                 'Germany':'Europe',
                 'India':'Asia',
                 'France':'Europe',
                 'South Korea':'Asia',
                 'Italy':'Europe',
                 'Spain':'Europe',
                 'Iran':'Asia',
                 'Australia':'Australia',
                 'Brazil':'South America'}
```

This function should return a DataFrame with index named Continent ['Asia', 'Australia', 'Europe', 'North America', 'South America'] and columns ['size', 'sum', 'mean', 'std']

```
In [38]: '''
@author: Steven Ponce
Date: 23 April 2021
'''

def answer_eleven():
    # YOUR CODE HERE
    #raise NotImplementedError()

    ContinentDict = {'China':'Asia',
                     'United States':'North America',
                     'Japan':'Asia',
                     'United Kingdom':'Europe',
                     'Russian Federation':'Europe',
                     'Canada':'North America',
                     'Germany':'Europe',
                     'India':'Asia',
                     'France':'Europe',
                     'South Korea':'Asia',
                     'Italy':'Europe',
                     'Spain':'Europe',
                     'Iran':'Asia',
                     'Australia':'Australia',
                     'Brazil':'South America'}

    df = answer_one()
    df['Population Estimate'] = df['Energy Supply'] / df['Energy Supply per Capita'].astype(float)

    stats = pd.DataFrame(columns = ['size', 'sum', 'mean', 'std'])

    for group, frame in df.groupby(ContinentDict):
        stats.loc[group] = [len(frame),
                           frame['Population Estimate'].sum(),
                           frame['Population Estimate'].mean(),
                           frame['Population Estimate'].std()]

    return stats
```

```
In [39]: answer_eleven()
```

Out[39]:

	size	sum	mean	std
Asia	5.0	2.898666e+09	5.797333e+08	6.790979e+08
Australia	1.0	2.331602e+07	2.331602e+07	NaN
Europe	6.0	4.579297e+08	7.632161e+07	3.464767e+07
North America	2.0	3.528552e+08	1.764276e+08	1.996696e+08
South America	1.0	2.059153e+08	2.059153e+08	NaN

```
In [40]: assert type(answer_eleven()) == pd.DataFrame, "Q11: You should return a DataFrame!"

assert answer_eleven().shape[0] == 5, "Q11: Wrong row numbers!"

assert answer_eleven().shape[1] == 4, "Q11: Wrong column numbers!"
```

Question 12

Cut % Renewable into 5 bins. Group Top15 by the Continent, as well as these new % Renewable bins. How many countries are in each of these groups?

This function should return a Series with a MultiIndex of Continent , then the bins for % Renewable . Do not include groups with no countries.

```
In [41]: '''
@author: Steven Ponce
Date: 23 April 2021
'''

def answer_twelve():
    # YOUR CODE HERE
    # raise NotImplementedError()

    ContinentDict = {'China':'Asia',
                     'United States':'North America',
                     'Japan':'Asia',
                     'United Kingdom':'Europe',
                     'Russian Federation':'Europe',
                     'Canada':'North America',
                     'Germany':'Europe',
                     'India':'Asia',
                     'France':'Europe',
                     'South Korea':'Asia',
                     'Italy':'Europe',
                     'Spain':'Europe',
                     'Iran':'Asia',
                     'Australia':'Australia',
                     'Brazil':'South America'}

    df = answer_one()
    df['Continent'] = pd.Series(ContinentDict)
    df['% Renewable']=pd.cut(df['% Renewable'],5)

    output = df.groupby(['Continent','% Renewable'])['Continent'].agg(np.size).dropna()

    return output
```

```
In [42]: answer_twelve()
```

```
Out[42]: Continent      % Renewable
Asia      (2.212, 15.753]      4.0
          (15.753, 29.227]      1.0
Australia (2.212, 15.753]      1.0
Europe    (2.212, 15.753]      1.0
          (15.753, 29.227]      3.0
          (29.227, 42.701]      2.0
North America (2.212, 15.753]      1.0
          (56.174, 69.648]      1.0
South America (56.174, 69.648]      1.0
Name: Continent, dtype: float64
```

```
In [43]: assert type(answer_twelve()) == pd.Series, "Q12: You should return a Series!"

assert len(answer_twelve()) == 9, "Q12: Wrong result numbers!"
```

Question 13

Convert the Population Estimate series to a string with thousands separator (using commas). Use all significant digits (do not round the results).

e.g. 12345678.90 -> 12,345,678.90

This function should return a series `PopEst` whose index is the country name and whose values are the population estimate string

```
In [44]: '''
@author: Steven Ponce
Date: 23 April 2021
'''

def answer_thirteen():
    # YOUR CODE HERE
    # raise NotImplementedError()

    df = answer_one()
    df['Population Estimate'] = df['Energy Supply'] / df['Energy Supply per Capita']
    PopEst = df['Population Estimate'].apply('{:,}'.format)

    return PopEst
```

```
In [45]: answer_thirteen()
```

```
Out[45]: Country
China      1,367,645,161.2903225
United States  317,615,384.61538464
Japan      127,409,395.97315437
United Kingdom  63,870,967.741935484
Russian Federation  143,500,000.0
Canada     35,239,864.86486486
Germany    80,369,696.96969697
India      1,276,730,769.2307692
France     63,837,349.39759036
South Korea  49,805,429.864253394
Italy      59,908,256.880733944
Spain      46,443,396.2264151
Iran       77,075,630.25210084
Australia  23,316,017.316017315
Brazil     205,915,254.23728815
Name: Population Estimate, dtype: object
```

```
In [46]: assert type(answer_thirteen()) == pd.Series, "Q13: You should return a Series!"
```

```
assert len(answer_thirteen()) == 15, "Q13: Wrong result numbers!"
```

Optional

Use the built in function `plot_optional()` to see an example visualization.


```
In [47]: def plot_optional():
import matplotlib as plt
%matplotlib inline
Top15 = answer_one()
ax = Top15.plot(x='Rank', y='% Renewable', kind='scatter',
                c=['#e41a1c', '#377eb8', '#e41a1c', '#4daf4a', '#4daf4a', '#377eb8', '#4daf4a', '#e41a1c',
                  '#4daf4a', '#e41a1c', '#4daf4a', '#4daf4a', '#e41a1c', '#dede00', '#ff7f00'],
                xticks=range(1,16), s=6*Top15['2014']/10**10, alpha=.75, figsize=[16,6]);

for i, txt in enumerate(Top15.index):
    ax.annotate(txt, [Top15['Rank'][i], Top15['% Renewable'][i]], ha='center')

print("This is an example of a visualization that can be created to help understand the data. \
This is a bubble chart showing % Renewable vs. Rank. The size of the bubble corresponds to the countries' \
2014 GDP, and the color corresponds to the continent.")
```

```
In [48]: plot_optional()
```

This is an example of a visualization that can be created to help understand the data. This is a bubble chart showing % Renewable vs. Rank. The size of the bubble corresponds to the countries' 2014 GDP, and the color corresponds to the continent.



```
In [ ]:
```