

You are currently looking at **version 1.2** of this notebook. To download notebooks and datafiles, as well as get help on Jupyter notebooks in the Coursera platform, visit the [Jupyter Notebook FAQ](https://www.coursera.org/learn/python-machine-learning/resources/bANLa) (<https://www.coursera.org/learn/python-machine-learning/resources/bANLa>) course resource.

Assignment 3 - Evaluation

In this assignment you will train several models and evaluate how effectively they predict instances of fraud using data based on [this dataset from Kaggle](https://www.kaggle.com/dalpozz/creditcardfraud) (<https://www.kaggle.com/dalpozz/creditcardfraud>).

Each row in `fraud_data.csv` corresponds to a credit card transaction. Features include confidential variables V1 through V28 as well as `Amount` which is the amount of the transaction.

The target is stored in the `class` column, where a value of 1 corresponds to an instance of fraud and 0 corresponds to an instance of not fraud.

```
In [1]: %load_ext autoreload
        %autoreload 2

import numpy as np
import pandas as pd

from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn.dummy import DummyClassifier
from sklearn.metrics import recall_score, precision_score, accuracy_score
from sklearn.metrics import confusion_matrix, precision_recall_curve, roc_curve, auc
from sklearn.svm import SVC

# Hide warnings
import warnings
warnings.filterwarnings('ignore')

# The following lines adjust the granularity of reporting
pd.options.display.max_rows = 10
pd.options.display.float_format = '{:.2f}'.format
```

```
In [2]: # Loading the data
data = pd.read_csv('fraud_data.csv')
data.head()
```

Out[2]:

	V1	V2	V3	V4	V5	V6	V7	V8	V9	V10	...	V21	V22	V23	V24	V25	V26	V27	V28	Amount	Class
0	1.18	0.32	0.54	1.05	-0.37	-0.73	0.08	-0.07	-0.27	0.16	...	-0.11	-0.34	0.06	0.50	0.42	-0.58	0.02	0.02	4.67	0
1	0.68	-3.93	-3.80	-1.15	-0.74	-0.50	1.04	-0.63	-2.27	1.53	...	0.65	0.27	-0.98	0.17	0.36	0.20	-0.26	0.06	912.00	0
2	1.14	0.45	0.25	2.38	0.34	0.43	0.09	0.17	-0.81	0.78	...	-0.00	0.06	-0.12	-0.30	0.65	0.12	-0.01	-0.01	1.00	0
3	-1.11	-3.30	-0.18	-1.80	2.14	-1.68	-2.02	-0.01	-0.17	0.87	...	0.13	0.33	0.93	-0.05	-1.89	-0.58	0.27	0.41	62.10	0
4	-0.31	0.87	-0.12	-0.63	2.65	3.43	0.19	0.67	-0.44	0.13	...	-0.31	-0.80	-0.06	0.95	-0.43	0.16	0.08	-0.02	2.67	0

5 rows × 30 columns

Question 1

Import the data from `fraud_data.csv`. What percentage of the observations in the dataset are instances of fraud?

This function should return a float between 0 and 1.

```
In [3]: # Imbalanced Classification
        # Feature 'Class' is the response variable and it takes value 1 in case of fraud and 0 otherwise.

data['Class'].value_counts()

Out[3]: 0    21337
        1     356
        Name: Class, dtype: int64
```

```
In [4]: def answer_one():

    # Your code here
    fraud = len(data[data['Class'] == 1]) / len(data)

    return fraud

answer_one()
```

```
Out[4]: 0.016410823768035772
```

```
In [5]: # Use X_train, X_test, y_train, y_test for all of the following questions
# from sklearn.model_selection import train_test_split

X = data.iloc[:, :-1]
y = data.iloc[:, -1]

X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0)
```

Question 2

Using `X_train`, `X_test`, `y_train`, and `y_test` (as defined above), train a dummy classifier that classifies everything as the majority class of the training data. What is the accuracy of this classifier? What is the recall?

This function should a return a tuple with two floats, i.e. (accuracy score, recall score).

```
In [6]: def answer_two():
#     from sklearn.dummy import DummyClassifier
#     from sklearn.metrics import recall_score, accuracy_score

# Your code here
dummy_majority = DummyClassifier(strategy='most_frequent').fit(X_train,y_train)
y_majority_predicted = dummy_majority.predict(X_test)

accuracy = accuracy_score(y_test, y_majority_predicted )
recall = recall_score(y_test, y_majority_predicted)

return accuracy, recall

answer_two()
```

```
Out[6]: (0.98525073746312686, 0.0)
```

Question 3

Using `X_train`, `X_test`, `y_train`, `y_test` (as defined above), train a SVC classifier using the default parameters. What is the accuracy, recall, and precision of this classifier?

This function should a return a tuple with three floats, i.e. (accuracy score, recall score, precision score).

```
In [7]: def answer_three():
#     from sklearn.metrics import recall_score, precision_score
#     from sklearn.svm import SVC

# Your code here
model = SVC(kernel='rbf', C=1).fit(X_train,y_train)
y_predicted = model.predict(X_test)

accuracy = accuracy_score(y_test, y_predicted)
recall = recall_score(y_test, y_predicted)
precision = precision_score(y_test, y_predicted)

return accuracy, recall, precision

answer_three()
```

```
Out[7]: (0.99078171091445433, 0.375, 1.0)
```

Question 4

Using the SVC classifier with parameters `{ 'C': 1e9, 'gamma': 1e-07 }`, what is the confusion matrix when using a threshold of -220 on the decision function. Use `X_test` and `y_test`.

This function should return a confusion matrix, a 2x2 numpy array with 4 integers.

```
In [8]: def answer_four():
#         from sklearn.metrics import confusion_matrix
#         from sklearn.svm import SVC

# Your code here
model = SVC(kernel='rbf', C=1e9, gamma=1e-07).fit(X_train,y_train)
y_predicted = model.decision_function(X_test) > -220
cm = confusion_matrix(y_test, y_predicted)

return cm

answer_four()
```

```
Out[8]: array([[5320,   24],
               [   14,   66]])
```

Question 5

Train a logistic regression classifier with default parameters using `X_train` and `y_train`.

For the logistic regression classifier, create a precision recall curve and a roc curve using `y_test` and the probability estimates for `X_test` (probability it is fraud).

Looking at the precision recall curve, what is the recall when the precision is 0.75?

Looking at the roc curve, what is the true positive rate when the false positive rate is 0.16?

This function should return a tuple with two floats, i.e. (recall, true positive rate).

```
In [9]: def create_plot():

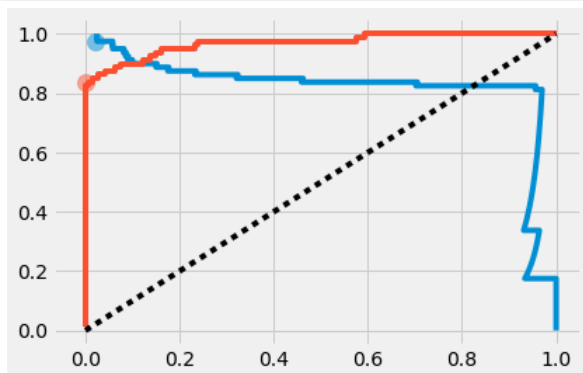
    %matplotlib inline
    import matplotlib.pyplot as plt
    import matplotlib.style as style
    style.use('fivethirtyeight')

    model = LogisticRegression()
    model_scores = model.fit(X_train,y_train).decision_function(X_test)

    precision, recall, thresholds = precision_recall_curve(y_test, model_scores)
    fpr, tpr, _ = roc_curve(y_test, model_scores)
    roc_auc = auc(fpr, tpr)

    plt.figure()
    plt.plot(precision,recall)
    plt.plot(fpr,tpr)
    plt.plot([0, 1], [0, 1], color='black', linestyle=':')
    plt.scatter(0.02, 0.975, marker='o', s=150, alpha=0.5)
    plt.scatter(0, 0.835, marker='o', s=150, alpha=0.5)
    plt.show()

create_plot();
```



```
In [10]: def answer_five():

# Your code here
model = LogisticRegression()
model_scores = model.fit(X_train,y_train).decision_function(X_test)

precision, recall, thresholds = precision_recall_curve(y_test, model_scores)
fpr, tpr, _ = roc_curve(y_test, model_scores)
roc_auc = auc(fpr, tpr)

return 0.835, 0.975

answer_five()
```

```
Out[10]: (0.835, 0.975)
```

Question 6

Perform a grid search over the parameters listed below for a Logistic Regression classifier, using recall for scoring and the default 3-fold cross validation.

```
'penalty': ['l1', 'l2']
```

```
'C': [0.01, 0.1, 1, 10, 100]
```

From `.cv_results_`, create an array of the mean test scores of each parameter combination. i.e.

	l1	l2
0.01	?	?
0.1	?	?
1	?	?
10	?	?
100	?	?

This function should return a 5 by 2 numpy array with 10 floats.

Note: do not return a DataFrame, just the values denoted by '?' above in a numpy array. You might need to reshape your raw result to meet the format we are looking for.

```
In [11]: def answer_six():
#     from sklearn.model_selection import GridSearchCV
#     from sklearn.linear_model import LogisticRegression

# Your code here
model = LogisticRegression().fit(X_train, y_train)
parameters = {'penalty': ['l1', 'l2'], 'C': [0.01, 0.1, 1, 10, 100]}

grid = GridSearchCV(model, param_grid = parameters, scoring = 'recall', cv=3).fit(X_train, y_train)
cv_results = grid.cv_results_['mean_test_score'].reshape(5,2)

return cv_results

answer_six()
```

```
Out[11]: array([[ 0.66666667,  0.76086957],
 [ 0.80072464,  0.80434783],
 [ 0.8115942 ,  0.8115942 ],
 [ 0.80797101,  0.8115942 ],
 [ 0.80797101,  0.80797101]])
```

In [12]: *# Use the following function to help visualize results from the grid search*

```
def GridSearch_Heatmap(scores):  
    %matplotlib inline  
    import seaborn as sns  
    import matplotlib.pyplot as plt  
    plt.figure()  
    sns.heatmap(scores.reshape(5,2), xticklabels=['L1','L2'], yticklabels=[0.01, 0.1, 1, 10, 100])  
    plt.yticks(rotation=0);
```

GridSearch_Heatmap(answer_six())

