
You are currently looking at **version 1.0** of this notebook. To download notebooks and datafiles, as well as get help on Jupyter notebooks in the Coursera platform, visit the [Jupyter Notebook FAQ \(https://www.coursera.org/learn/python-text-mining/resources/d9pwm\)](https://www.coursera.org/learn/python-text-mining/resources/d9pwm) course resource.

Assignment 2 - Introduction to NLTK

In part 1 of this assignment you will use nltk to explore the Herman Melville novel Moby Dick. Then in part 2 you will create a spelling recommender function that uses nltk to find words similar to the misspelling.

Part 1 - Analyzing Moby Dick

```
In [1]: import nltk
nltk.download("book")
from nltk.book import *

# text1: Moby Dick by Herman Melville 1851
```

```
[nltk_data] Downloading collection 'book'
[nltk_data] |
[nltk_data] | Downloading package abc to /home/jovyan/nltk_data...
[nltk_data] | Package abc is already up-to-date!
[nltk_data] | Downloading package brown to
[nltk_data] | /home/jovyan/nltk_data...
[nltk_data] | Package brown is already up-to-date!
[nltk_data] | Downloading package chat80 to
[nltk_data] | /home/jovyan/nltk_data...
[nltk_data] | Package chat80 is already up-to-date!
[nltk_data] | Downloading package cmudict to
[nltk_data] | /home/jovyan/nltk_data...
[nltk_data] | Package cmudict is already up-to-date!
[nltk_data] | Downloading package conll2000 to
[nltk_data] | /home/jovyan/nltk_data...
[nltk_data] | Package conll2000 is already up-to-date!
[nltk_data] | Downloading package conll2002 to
[nltk_data] | /home/jovyan/nltk_data...
[nltk_data] | Package conll2002 is already up-to-date!
[nltk_data] | Downloading package dependency_treebank to
[nltk_data] | /home/jovyan/nltk_data...
[nltk_data] | Package dependency_treebank is already up-to-date!
[nltk_data] | Downloading package genesis to
[nltk_data] | /home/jovyan/nltk_data...
[nltk_data] | Package genesis is already up-to-date!
[nltk_data] | Downloading package gutenberg to
[nltk_data] | /home/jovyan/nltk_data...
[nltk_data] | Package gutenberg is already up-to-date!
[nltk_data] | Downloading package ier to /home/jovyan/nltk_data...
[nltk_data] | Package ier is already up-to-date!
[nltk_data] | Downloading package inaugural to
[nltk_data] | /home/jovyan/nltk_data...
[nltk_data] | Package inaugural is already up-to-date!
[nltk_data] | Downloading package movie_reviews to
[nltk_data] | /home/jovyan/nltk_data...
[nltk_data] | Package movie_reviews is already up-to-date!
[nltk_data] | Downloading package nps_chat to
[nltk_data] | /home/jovyan/nltk_data...
[nltk_data] | Package nps_chat is already up-to-date!
[nltk_data] | Downloading package names to
[nltk_data] | /home/jovyan/nltk_data...
[nltk_data] | Package names is already up-to-date!
[nltk_data] | Downloading package ppattach to
[nltk_data] | /home/jovyan/nltk_data...
[nltk_data] | Package ppattach is already up-to-date!
[nltk_data] | Downloading package reuters to
[nltk_data] | /home/jovyan/nltk_data...
[nltk_data] | Package reuters is already up-to-date!
[nltk_data] | Downloading package senseval to
[nltk_data] | /home/jovyan/nltk_data...
[nltk_data] | Package senseval is already up-to-date!
[nltk_data] | Downloading package state_union to
[nltk_data] | /home/jovyan/nltk_data...
[nltk_data] | Package state_union is already up-to-date!
[nltk_data] | Downloading package stopwords to
[nltk_data] | /home/jovyan/nltk_data...
[nltk_data] | Package stopwords is already up-to-date!
[nltk_data] | Downloading package swadesh to
[nltk_data] | /home/jovyan/nltk_data...
[nltk_data] | Package swadesh is already up-to-date!
[nltk_data] | Downloading package timit to
[nltk_data] | /home/jovyan/nltk_data...
[nltk_data] | Package timit is already up-to-date!
[nltk_data] | Downloading package treebank to
[nltk_data] | /home/jovyan/nltk_data...
[nltk_data] | Package treebank is already up-to-date!
[nltk_data] | Downloading package toolbox to
[nltk_data] | /home/jovyan/nltk_data...
[nltk_data] | Package toolbox is already up-to-date!
[nltk_data] | Downloading package udhr to /home/jovyan/nltk_data...
[nltk_data] | Package udhr is already up-to-date!
[nltk_data] | Downloading package udhr2 to
[nltk_data] | /home/jovyan/nltk_data...
[nltk_data] | Package udhr2 is already up-to-date!
[nltk_data] | Downloading package unicode_samples to
[nltk_data] | /home/jovyan/nltk_data...
```

```

[nltk_data] | Package unicode_samples is already up-to-date!
[nltk_data] | Downloading package webtext to
[nltk_data] | /home/jovyan/nltk_data...
[nltk_data] | Package webtext is already up-to-date!
[nltk_data] | Downloading package wordnet to
[nltk_data] | /home/jovyan/nltk_data...
[nltk_data] | Package wordnet is already up-to-date!
[nltk_data] | Downloading package wordnet_ic to
[nltk_data] | /home/jovyan/nltk_data...
[nltk_data] | Package wordnet_ic is already up-to-date!
[nltk_data] | Downloading package words to
[nltk_data] | /home/jovyan/nltk_data...
[nltk_data] | Package words is already up-to-date!
[nltk_data] | Downloading package maxent_treebank_pos_tagger to
[nltk_data] | /home/jovyan/nltk_data...
[nltk_data] | Package maxent_treebank_pos_tagger is already up-
[nltk_data] | to-date!
[nltk_data] | Downloading package maxent_ne_chunker to
[nltk_data] | /home/jovyan/nltk_data...
[nltk_data] | Package maxent_ne_chunker is already up-to-date!
[nltk_data] | Downloading package universal_tagset to
[nltk_data] | /home/jovyan/nltk_data...
[nltk_data] | Package universal_tagset is already up-to-date!
[nltk_data] | Downloading package punkt to
[nltk_data] | /home/jovyan/nltk_data...
[nltk_data] | Package punkt is already up-to-date!
[nltk_data] | Downloading package book_grammars to
[nltk_data] | /home/jovyan/nltk_data...
[nltk_data] | Package book_grammars is already up-to-date!
[nltk_data] | Downloading package city_database to
[nltk_data] | /home/jovyan/nltk_data...
[nltk_data] | Package city_database is already up-to-date!
[nltk_data] | Downloading package tagsets to
[nltk_data] | /home/jovyan/nltk_data...
[nltk_data] | Package tagsets is already up-to-date!
[nltk_data] | Downloading package panlex_swadesh to
[nltk_data] | /home/jovyan/nltk_data...
[nltk_data] | Package panlex_swadesh is already up-to-date!
[nltk_data] | Downloading package averaged_perceptron_tagger to
[nltk_data] | /home/jovyan/nltk_data...
[nltk_data] | Package averaged_perceptron_tagger is already up-
[nltk_data] | to-date!
[nltk_data] |
[nltk_data] | Done downloading collection book
*** Introductory Examples for the NLTK Book ***
Loading text1, ..., text9 and sent1, ..., sent9
Type the name of the text or sentence to view it.
Type: 'texts()' or 'sents()' to list the materials.
text1: Moby Dick by Herman Melville 1851
text2: Sense and Sensibility by Jane Austen 1811
text3: The Book of Genesis
text4: Inaugural Address Corpus
text5: Chat Corpus
text6: Monty Python and the Holy Grail
text7: Wall Street Journal
text8: Personals Corpus
text9: The Man Who Was Thursday by G . K . Chesterton 1908

```

```

In [2]: import nltk, collections
import pandas as pd
import numpy as np

# If you would like to work with the raw text you can use 'moby_raw'
with open('moby.txt', 'r') as f:
    moby_raw = f.read()

# If you would like to work with the novel in nltk.Text format you can use 'text1'
moby_tokens = nltk.word_tokenize(moby_raw)
text1 = nltk.Text(moby_tokens)

```

Example 1

How many tokens (words and punctuation symbols) are in text1?

This function should return an integer.

```
In [3]: def example_one():  
        return len(nltk.word_tokenize(moby_raw)) # or alternatively len(text1)  
        example_one()
```

Out[3]: 254989

Example 2

How many unique tokens (unique words and punctuation) does text1 have?

This function should return an integer.

```
In [4]: def example_two():  
        return len(set(nltk.word_tokenize(moby_raw))) # or alternatively len(set(text1))  
        example_two()
```

Out[4]: 20755

Example 3

After lemmatizing the verbs, how many unique tokens does text1 have?

This function should return an integer.

```
In [5]: from nltk.stem import WordNetLemmatizer  
  
def example_three():  
    lemmatizer = WordNetLemmatizer()  
    lemmatized = [lemmatizer.lemmatize(w,'v') for w in text1]  
  
    return len(set(lemmatized))  
  
example_three()
```

Out[5]: 16900

Question 1

What is the lexical diversity of the given text input? (i.e. ratio of unique tokens to the total number of tokens)

This function should return a float.

```
In [6]: def answer_one():  
        result = example_two() / example_one()  
  
        return result  
  
answer_one()
```

Out[6]: 0.08139566804842562

Question 2

What percentage of tokens is 'whale' or 'Whale'?

This function should return a float.

```
In [7]: def answer_two():  
  
        dist = nltk.FreqDist(text1)  
  
        result = (dist['whale'] + dist['Whale']) / example_one()*100  
  
        return result  
  
answer_two()
```

Out[7]: 0.4125668166077752

Question 3

What are the 20 most frequently occurring (unique) tokens in the text? What is their frequency?

This function should return a list of 20 tuples where each tuple is of the form (token, frequency). The list should be sorted in descending order of frequency.

```
In [8]: def answer_three():  
  
        result = nltk.FreqDist(text1).most_common(20)  
  
        return result  
  
answer_three()
```

```
Out[8]: [(' ', 19204),  
        ('the', 13715),  
        ('.', 7308),  
        ('of', 6513),  
        ('and', 6010),  
        ('a', 4545),  
        ('to', 4515),  
        (';', 4173),  
        ('in', 3908),  
        ('that', 2978),  
        ('his', 2459),  
        ('it', 2196),  
        ('I', 2097),  
        ('!', 1767),  
        ('is', 1722),  
        ('--', 1713),  
        ('with', 1659),  
        ('he', 1658),  
        ('was', 1639),  
        ('as', 1620)]
```

Question 4

What tokens have a length of greater than 5 and frequency of more than 150?

This function should return an alphabetically sorted list of the tokens that match the above constraints. To sort your list, use sorted()

```
In [9]: def answer_four():

    dist = nltk.FreqDist(text1)

    vocabulary = dist.keys()

    words = [w for w in vocabulary if len(w) > 5 and dist[w] > 150]

    result = sorted(words)

    return result

answer_four()
```

```
Out[9]: ['Captain',
        'Pequod',
        'Queequeg',
        'Starbuck',
        'almost',
        'before',
        'himself',
        'little',
        'seemed',
        'should',
        'though',
        'through',
        'whales',
        'without']
```

Question 5

Find the longest word in text1 and that word's length.

This function should return a tuple (Longest_word, Length).

```
In [10]: def answer_five():

    word_length = max(len(w) for w in text1)

    word = [w for w in text1 if len(w) == word_length]

    result = word[0], word_length

    return result

answer_five()
```

```
Out[10]: ("twelve-o'clock-at-night", 23)
```

Question 6

What unique words have a frequency of more than 2000? What is their frequency?

"Hint: you may want to use `isalpha()` to check if the token is a word and not punctuation."

This function should return a list of tuples of the form (frequency, word) sorted in descending order of frequency.

```
In [11]: def answer_six():

    dist = nltk.FreqDist(text1)

    words = [(dist[w],w) for w in set(text1) if w.isalpha() and dist[w] > 2000]

    result = sorted(words, reverse=True)

    return result

answer_six()
```

```
Out[11]: [(13715, 'the'),
(6513, 'of'),
(6010, 'and'),
(4545, 'a'),
(4515, 'to'),
(3908, 'in'),
(2978, 'that'),
(2459, 'his'),
(2196, 'it'),
(2097, 'I')]
```

Question 7

What is the average number of tokens per sentence?

This function should return a float.

```
In [12]: def answer_seven():

    result = np.mean([len(nltk.word_tokenize(sent)) for sent in nltk.sent_tokenize(moby_raw)])

    return result

answer_seven()
```

```
Out[12]: 25.881952902963864
```

Question 8

What are the 5 most frequent parts of speech in this text? What is their frequency?

This function should return a list of tuples of the form (part_of_speech, frequency) sorted in descending order of frequency.

```
In [13]: def answer_eight():

    tags = nltk.pos_tag(text1)

    counts = collections.Counter([x[1] for x in tags])

    result = counts.most_common(5)

    return result

answer_eight()
```

```
Out[13]: [('NN', 32730), ('IN', 28657), ('DT', 25867), ('', 19204), ('JJ', 17620)]
```


Part 2 - Spelling Recommender

For this part of the assignment you will create three different spelling recommenders, that each take a list of misspelled words and recommends a correctly spelled word for every word in the list.

For every misspelled word, the recommender should find the word in `correct_spellings` that has the shortest distance*, and starts with the same letter as the misspelled word, and return that word as a recommendation.

*Each of the three different recommenders will use a different distance measure (outlined below).

Each of the recommenders should provide recommendations for the three default words provided: ['cormulent', 'incendenece', 'validate'].

```
In [14]: from nltk.corpus import words

correct_spellings = words.words()
```

Question 9

For this recommender, your function should provide recommendations for the three default words provided above using the following distance metric:

Jaccard distance (https://en.wikipedia.org/wiki/Jaccard_index) on the trigrams of the two words.

This function should return a list of length three: ['cormulent_reccommendation', 'incendenece_reccommendation', 'validate_reccommendation'].

```
In [15]: def answer_nine(entries=['cormulent', 'incendenece', 'validate']):

    result = []

    for entry in entries:

        spelling = [x for x in correct_spellings if x[0] == entry[0] and len(x) >= 5]

        Jaccard_distance = [nltk.jaccard_distance(set(nltk.ngrams(entry,n=3)), set(nltk.ngrams(x,n=3))) for x
in spelling]

        result.append(spelling[np.argmin(Jaccard_distance)])

    return result

answer_nine()
```

```
Out[15]: ['corpulent', 'indecence', 'validate']
```

Question 10

For this recommender, your function should provide recommendations for the three default words provided above using the following distance metric:

Jaccard distance (https://en.wikipedia.org/wiki/Jaccard_index) on the 4-grams of the two words.

This function should return a list of length three: ['cormulent_reccommendation', 'incendenece_reccommendation', 'validate_reccommendation'].

```
In [16]: def answer_ten(entries=['cormulent', 'incendenece', 'validate']):
    result = []
    for entry in entries:
        spelling = [x for x in correct_spellings if x[0] == entry[0] and len(x) >= 5]
        Jaccard_distance = [nltk.jaccard_distance(set(nltk.ngrams(entry,n=4)), set(nltk.ngrams(x,n=4))) for x
in spelling]
        result.append(spelling[np.argmin(Jaccard_distance)])
    return result
answer_ten()
```

Out[16]: ['cormus', 'incendiary', 'valid']

Question 11

For this recommender, your function should provide recommendations for the three default words provided above using the following distance metric:

Edit distance on the two words with transpositions. (https://en.wikipedia.org/wiki/Damerau%E2%80%93Levenshtein_distance)

This function should return a list of length three: ['cormulent_reccomendation', 'incendenece_reccomendation', 'validate_reccomendation'].

```
In [17]: def answer_eleven(entries=['cormulent', 'incendenece', 'validate']):
    result = []
    for entry in entries:
        spelling = [x for x in correct_spellings if x[0] == entry[0] and len(x) >= 5]
        Damerau_Levenshtein_distance = [nltk.edit_distance(x, entry, transpositions=True) for x in spelling]
        result.append(spelling[np.argmin(Damerau_Levenshtein_distance)])
    return result
answer_eleven()
```

Out[17]: ['corpulent', 'intendence', 'validate']